


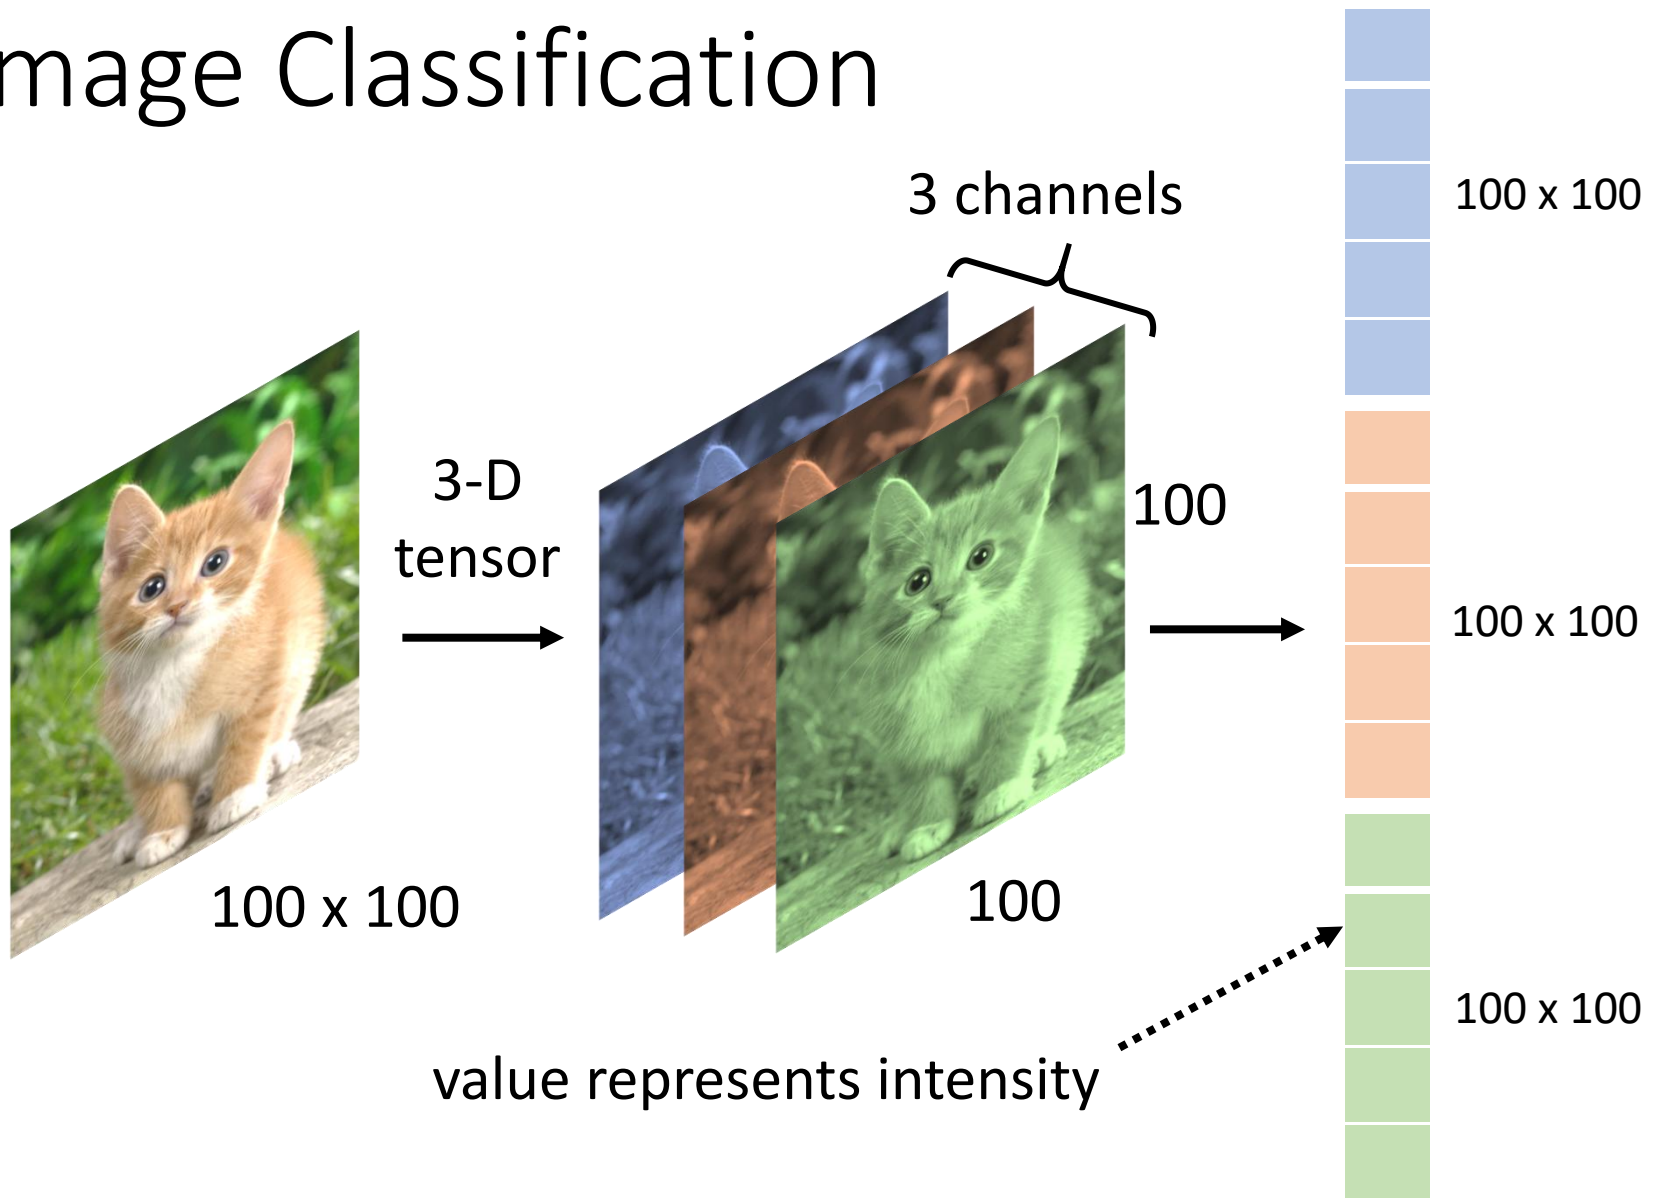
Convolution Neural Network (CNN)

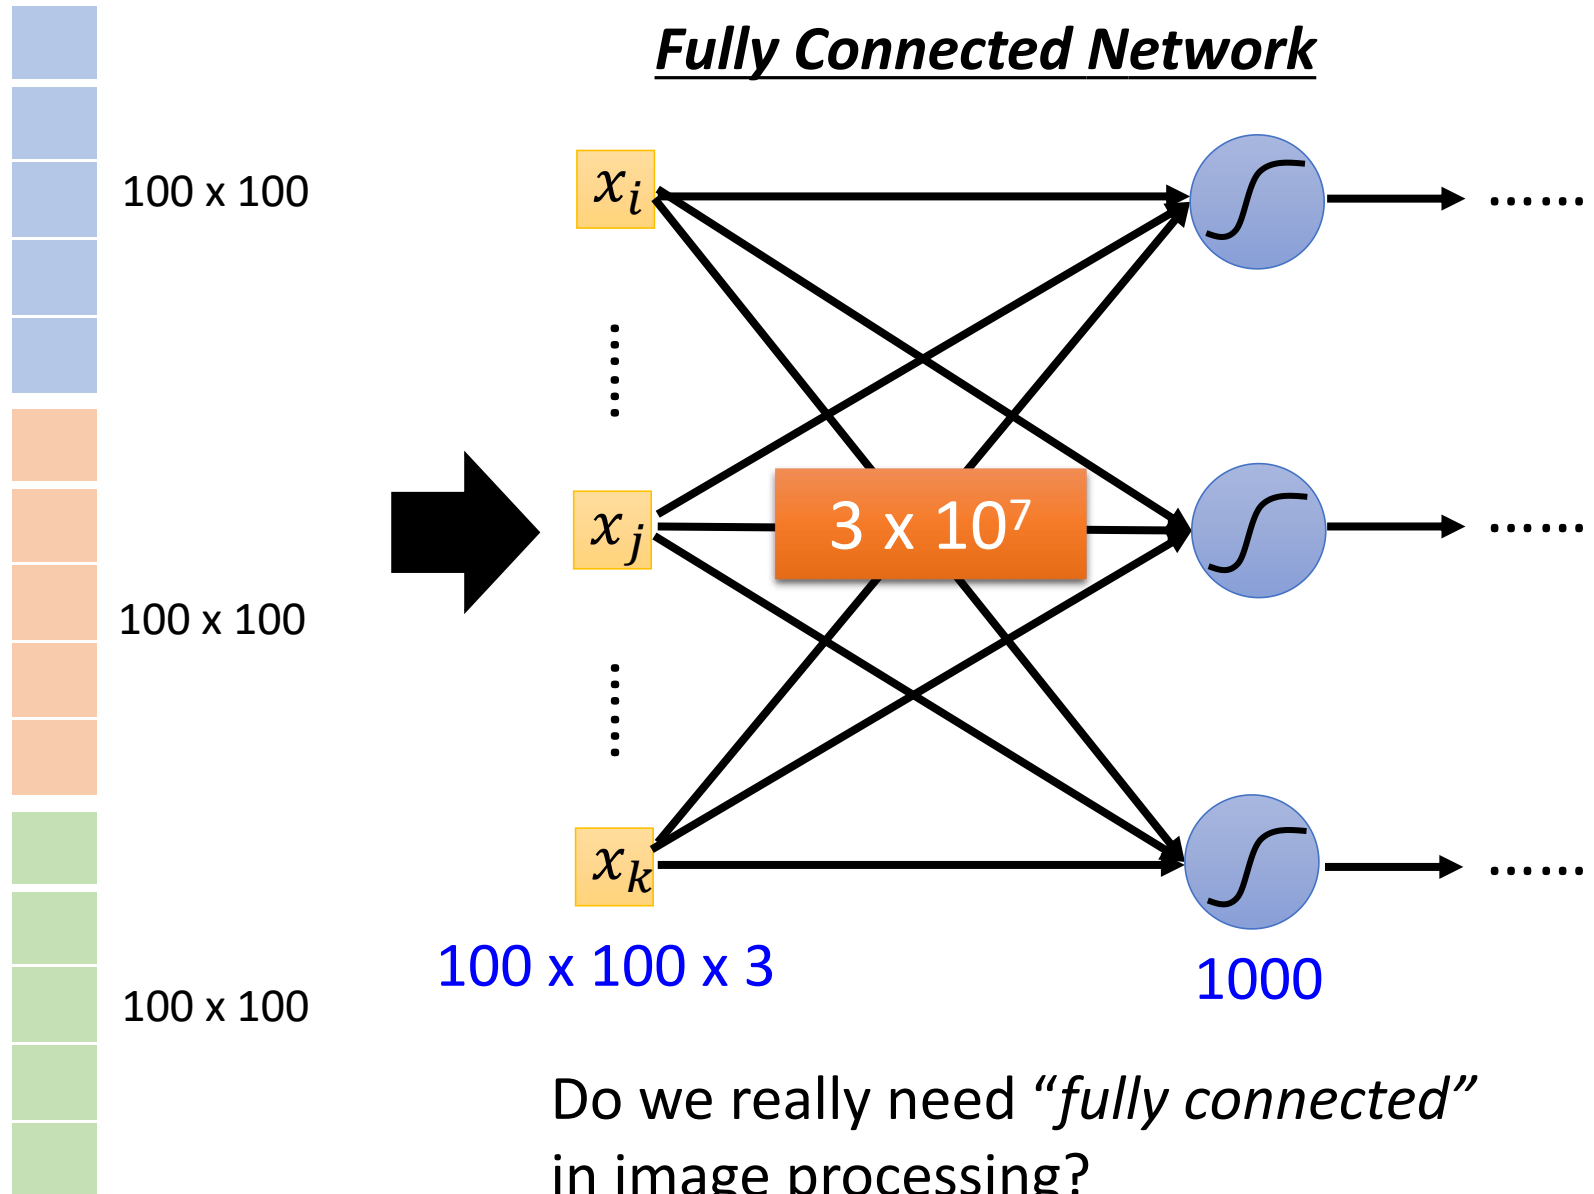
Image Classification



(All the images to be classified have the same size.)

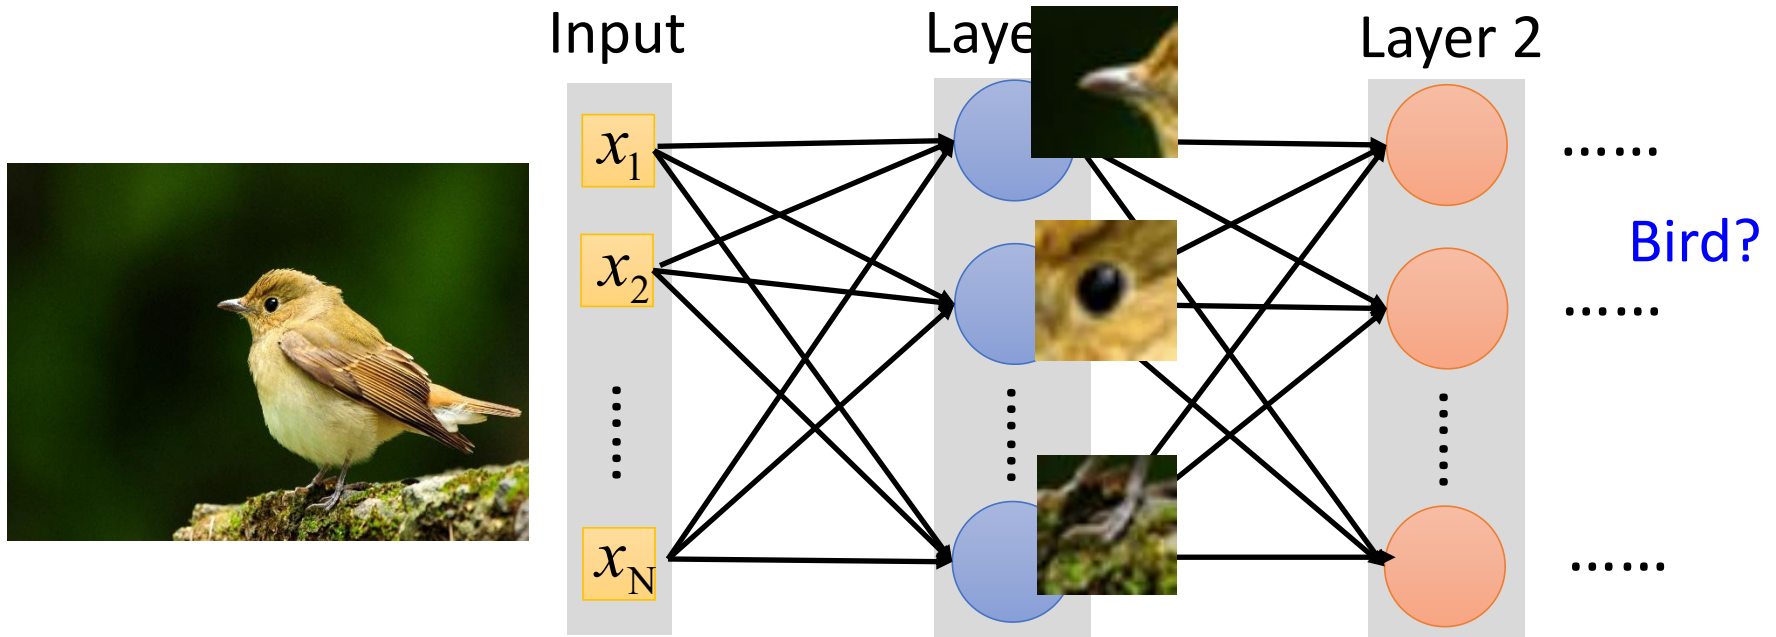
Image Classification





Observation 1

Identifying some critical patterns




Perhaps human also identify birds in a similar way ... 😊

Observation 1

A neuron does not have to see the whole image.

Need to see the whole image?



Input




Layer 1



basic
detector

Layer 2





advanced
detector

.....


bird
.....

.....



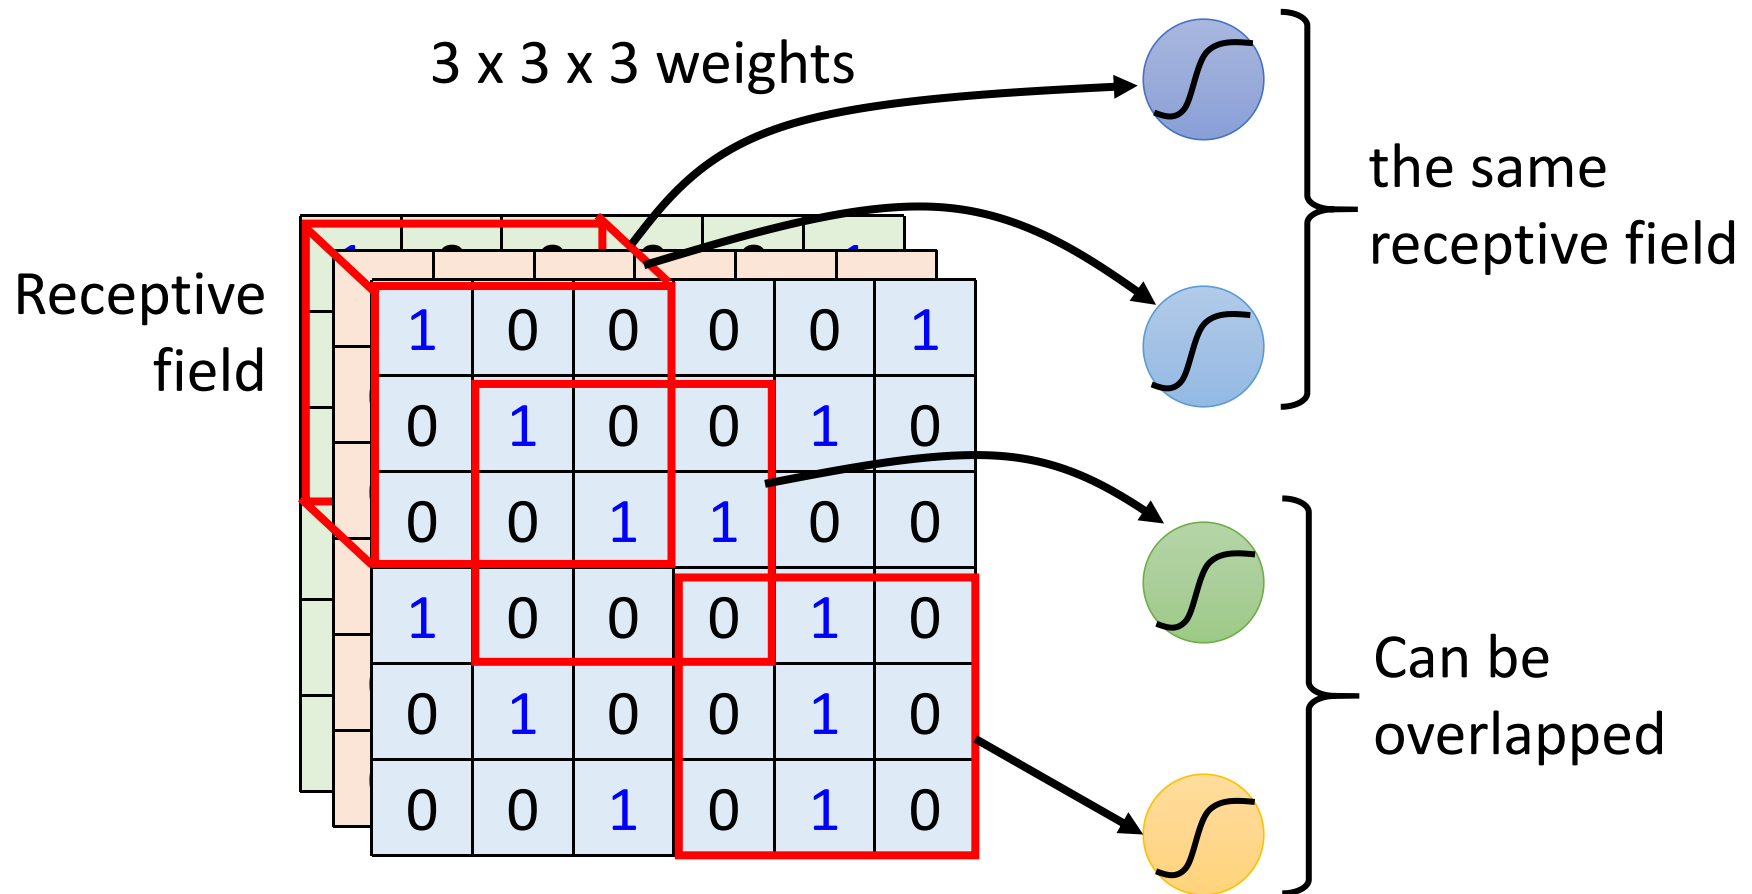
Some patterns are much smaller than the whole image.

Simplification 1



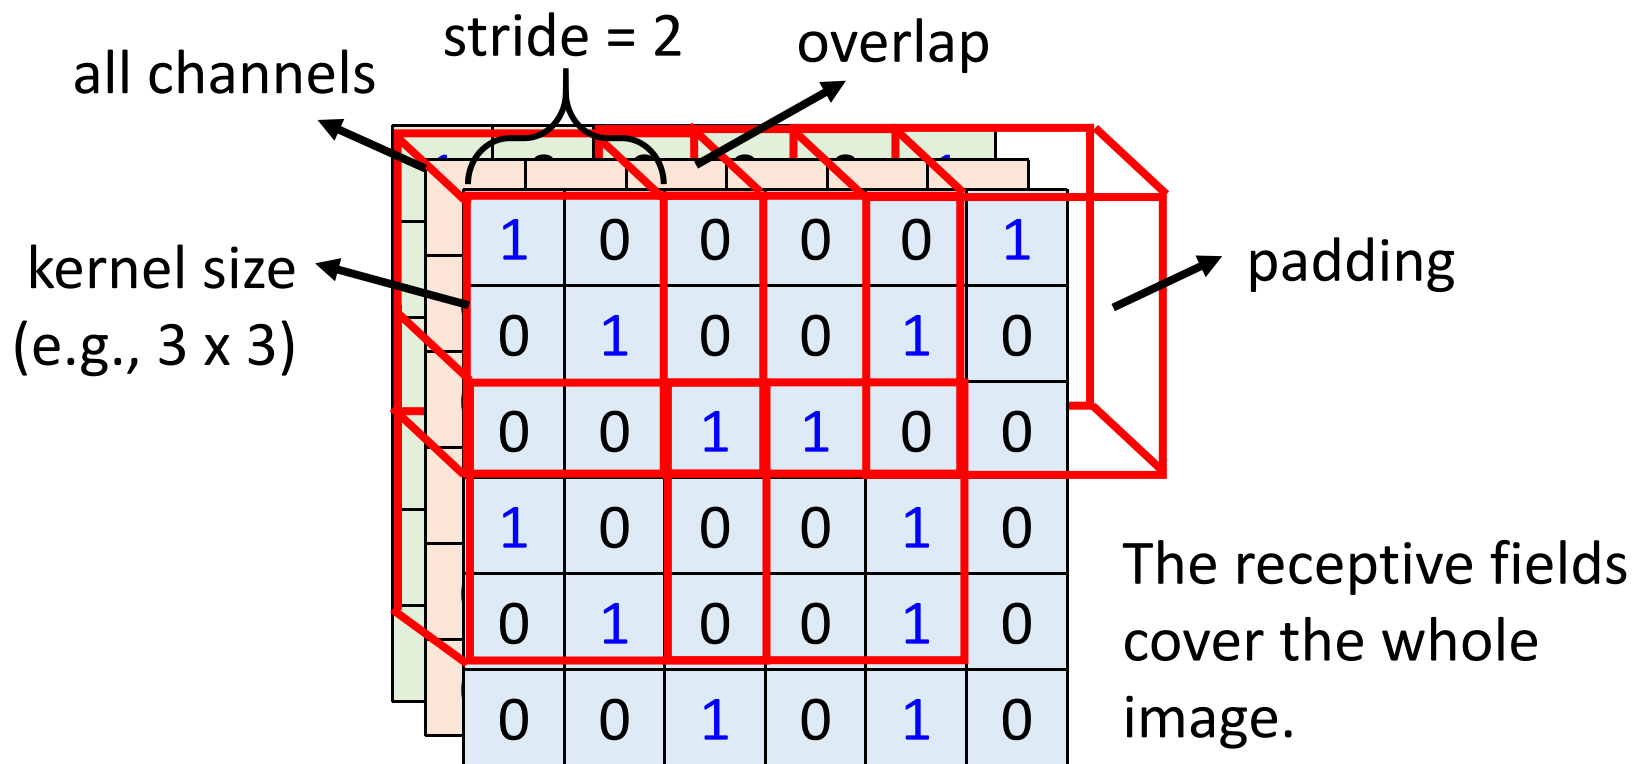
Simplification 1

- Can different neurons have different sizes of receptive field?
- Cover only some channels?
- Not square receptive field?



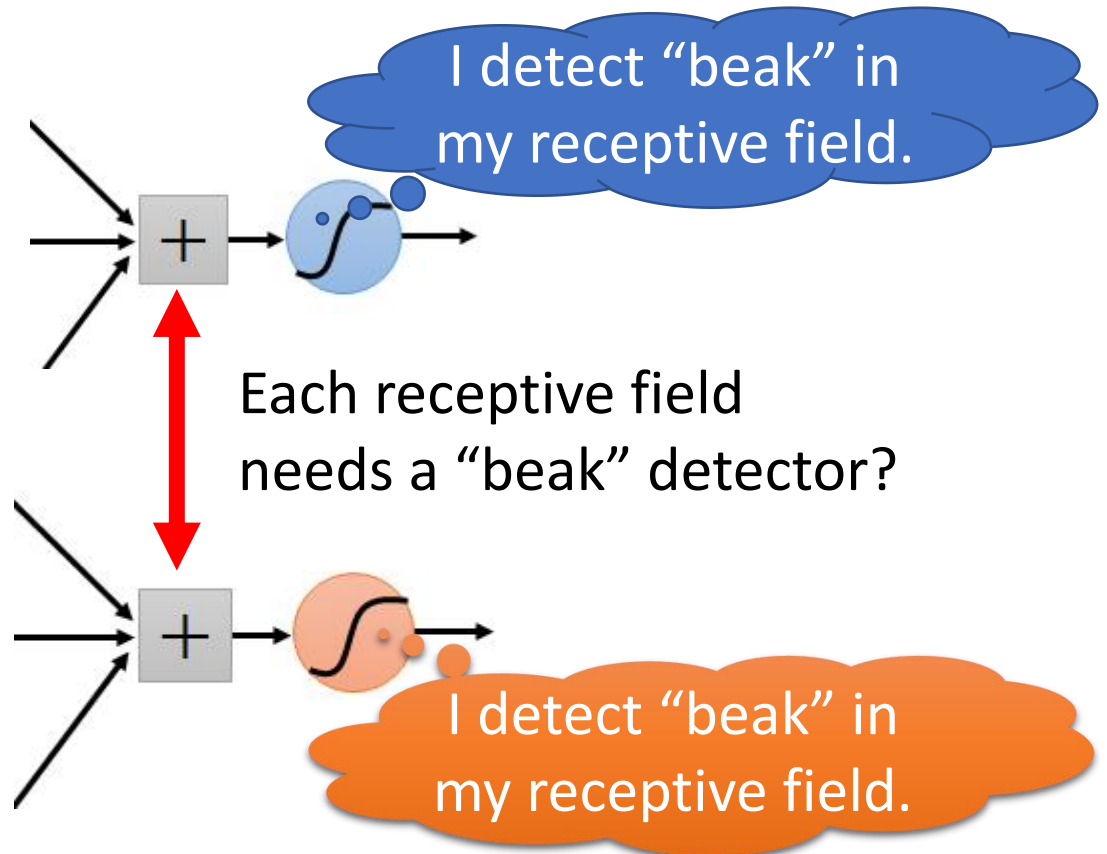
Simplification 1 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).





Observation 2


- The same patterns appear in different regions.





Simplification 2



Simplification 2




Two neurons with the same receptive field would not share parameters.



Simplification 2 – Typical Setting


Each receptive field has a set of neurons (e.g., 64 neurons).



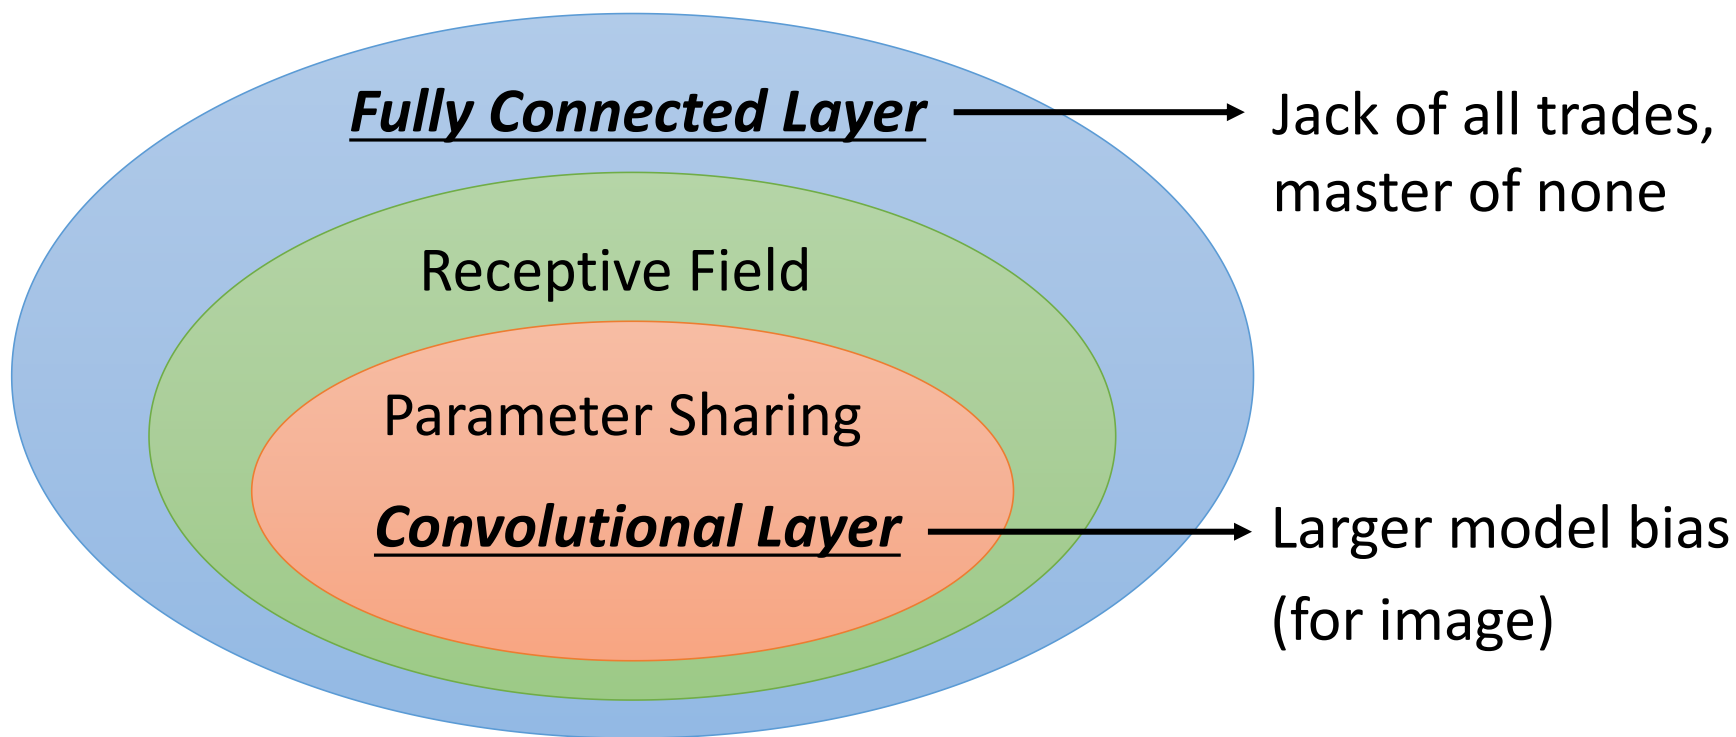
Simplification 2 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).

Each receptive field has the neurons with the same set of parameters.



Benefit of Convolutional Layer



- Some patterns are much smaller than the whole image.
- The same patterns appear in different regions.

Convolutional Layer




Convolution




...

channel = 3 (colorful)

channel = 1 (black and white)



Filter 1
3 x 3 x channel
tensor



Filter 2
3 x 3 x channel
tensor

...

Each filter detects a small pattern (3 x 3 x channel).

Convolutional Layer

Consider channel = 1
(black and white image)

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮

(The values in the filters
are unknown parameters.)

Convolutional Layer

stride=1

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Convolutional Layer

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0


6 x 6 image

Do the same process for every filter



Quiz

- Given a $N \times N$ size picture, filter size $F \times F$, stride S , what is the size of output feature map? (No padding)




CSDN @Rainbowman 0

Quiz

- Given a $N \times N$ size picture, filter size $F \times F$, stride S , what is the size of output feature map?

$$outputsize = \frac{N - F}{S} + 1$$




CSDN @Rainbowman 0

Quiz

- Given a $N \times N$ size picture, filter size $F \times F$, stride S , padding size P , what is the size of output feature map?

$$(N-F+2P)/S+1$$



Convolutional Layer




64
filters

Convolution

Convolution

⋮



"Image" with 64 channels

Multiple Convolutional Layers




64
filters


Convolution

Convolution

⋮



"Image" with 64 channels



Multiple Convolutional Layers




64
filters

Convolution

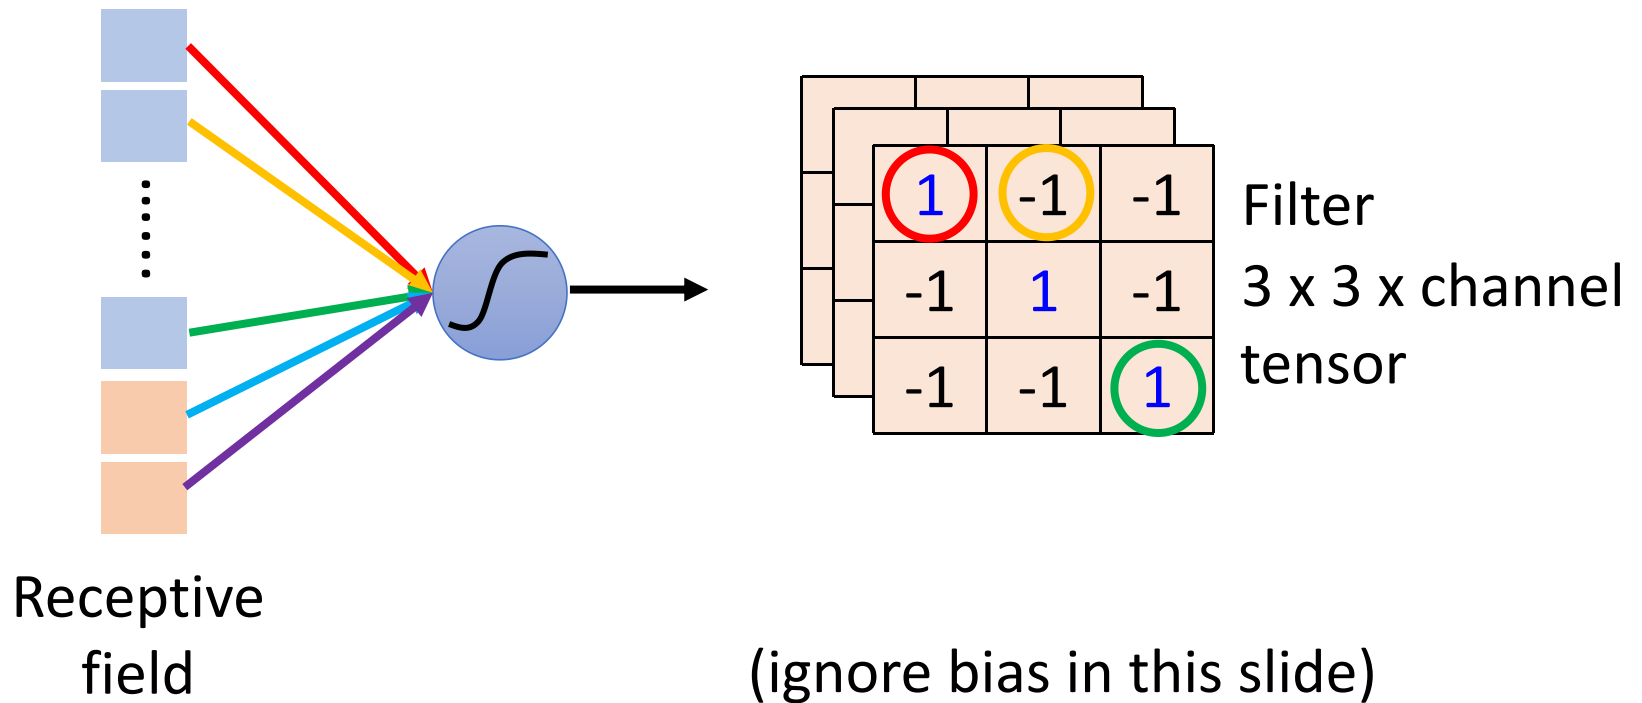
Convolution

⋮


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0




Comparison of Two Stories



The neurons with different receptive fields **share the parameters**.



Each filter convolves over the input image.



Convolutional Layer

<u>Neuron Version Story</u>	<u>Filter Version Story</u>
Each neuron only considers a receptive field.	There are a set of filters detecting small patterns.
The neurons with different receptive fields share the parameters.	Each filter convolves over the input image.

They are the same story.

Observation 3

- Subsampling the pixels will not change the object

bird



subsampling

bird



Pooling – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

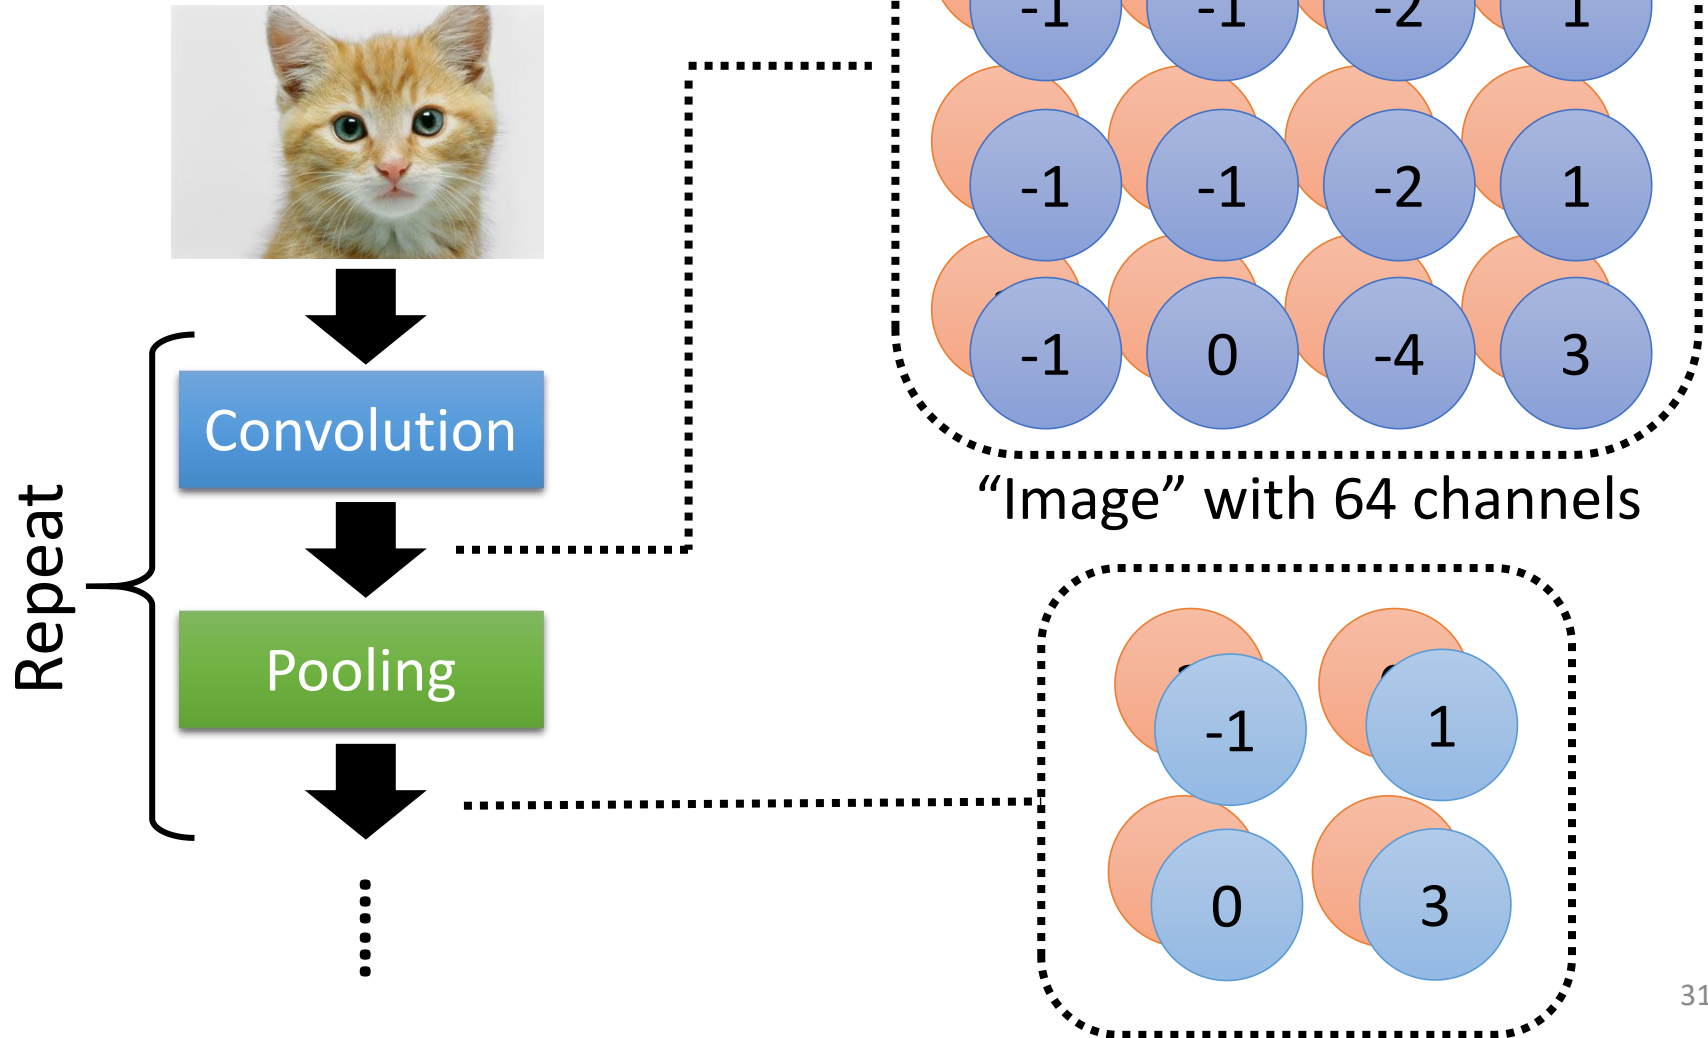
-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

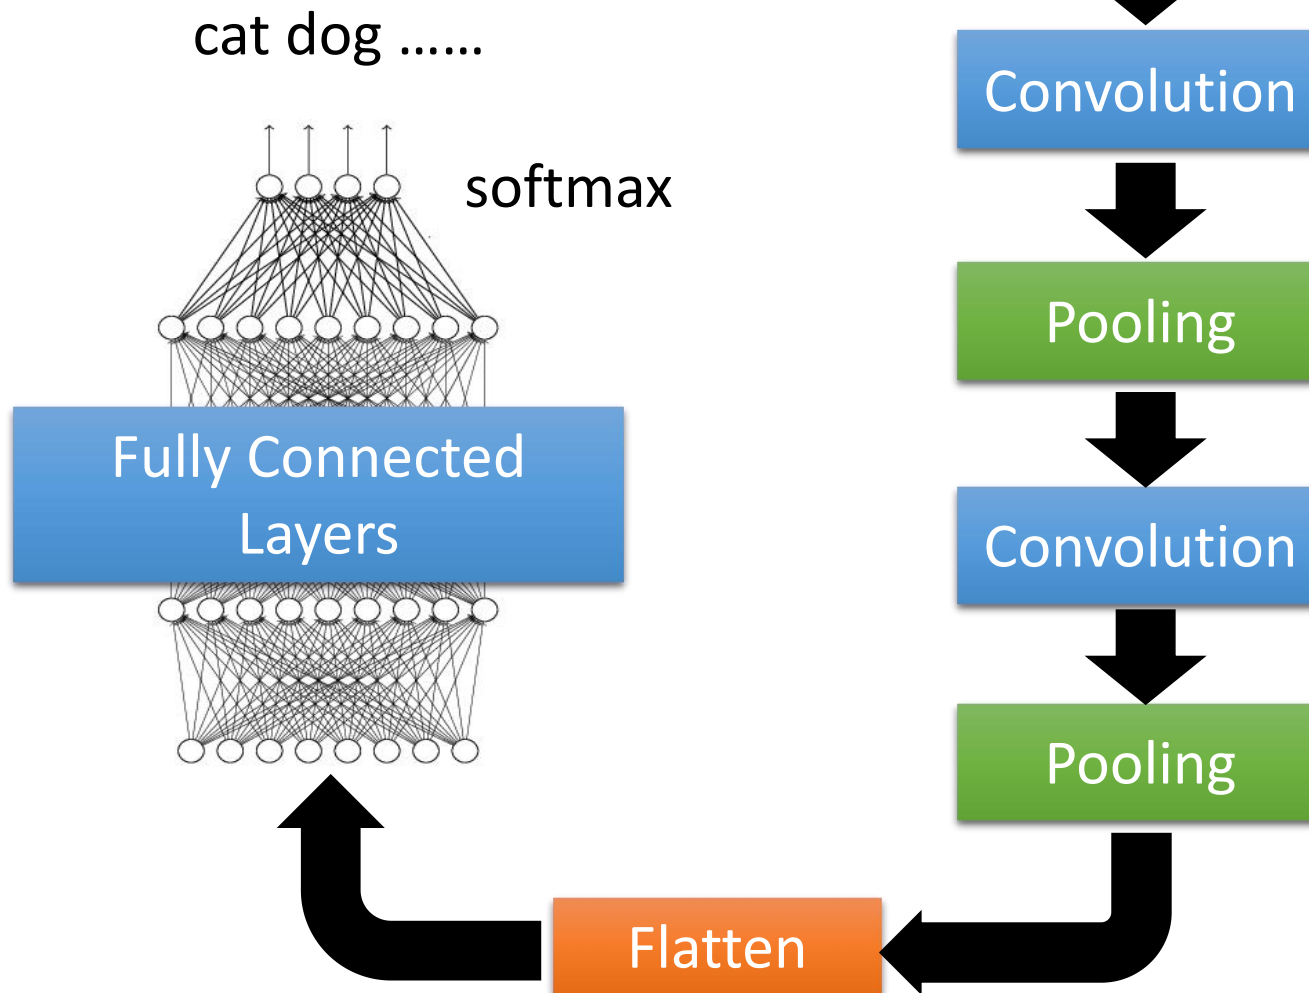
3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3

Convolutional Layers + Pooling



The whole CNN



Application: Playing Go



48 channels
in Alpha Go

Black: 1
white: -1
none: 0



Next move
(19 x 19
positions)
19 x 19 classes


Fully-connected
network can be used

But CNN performs much better.



Why CNN for Go playing?

- Some patterns are much smaller than the whole image

Alpha Go uses 5 x 5 for first layer



- The same patterns appear in different regions.



Why CNN for Go playing?

- Subsampling the pixels will not change the object



Pooling

How to explain this???

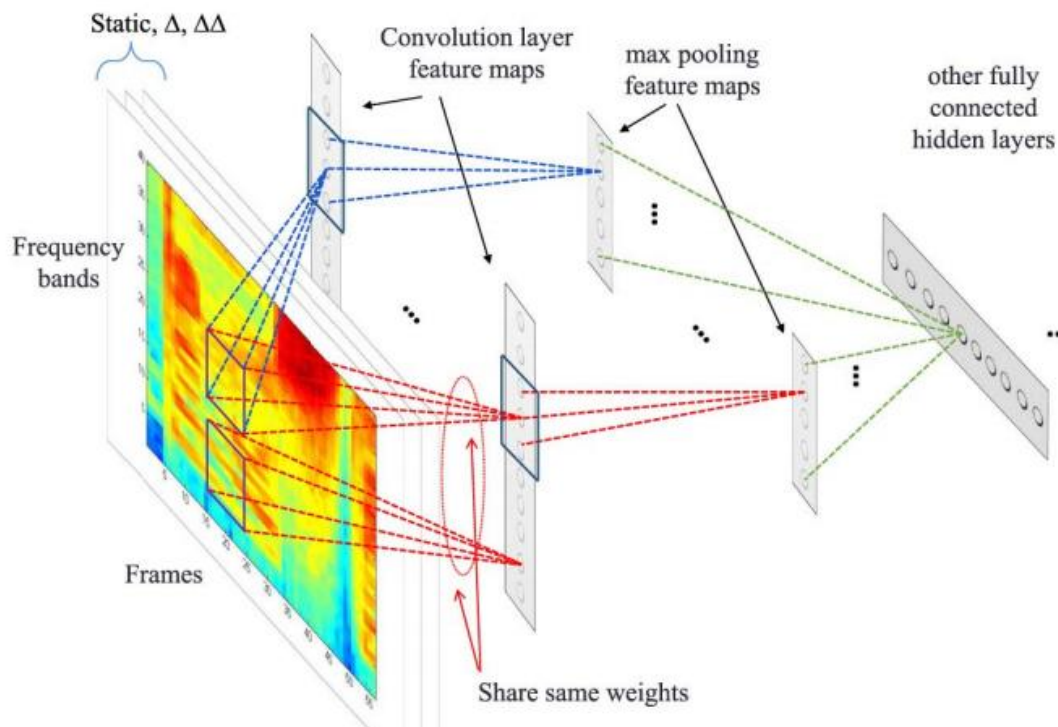
Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used $k = 192$ filters; Fig. 2b and Extended Data Table 1. The match version of AlphaGo used 256 and 384 filters.

Alpha Go does not use Pooling

More Applications

Speech

<https://dl.acm.org/doi/10.1109/TASLP.2014.2339736>



Natural Language Processing

<https://www.aclweb.org/anthology/S15-2079/>

